

# External Sorting

- Adapt fastest internal-sort methods.
- ✓ Quick sort ...best average run time.
- Merge sort ... best worst-case run time.

# Internal Merge Sort Review

- Phase 1
  - Create initial sorted segments
    - Natural segments
    - Insertion sort
- Phase 2
  - Merge pairs of sorted segments, in merge passes, until only **1** segment remains.

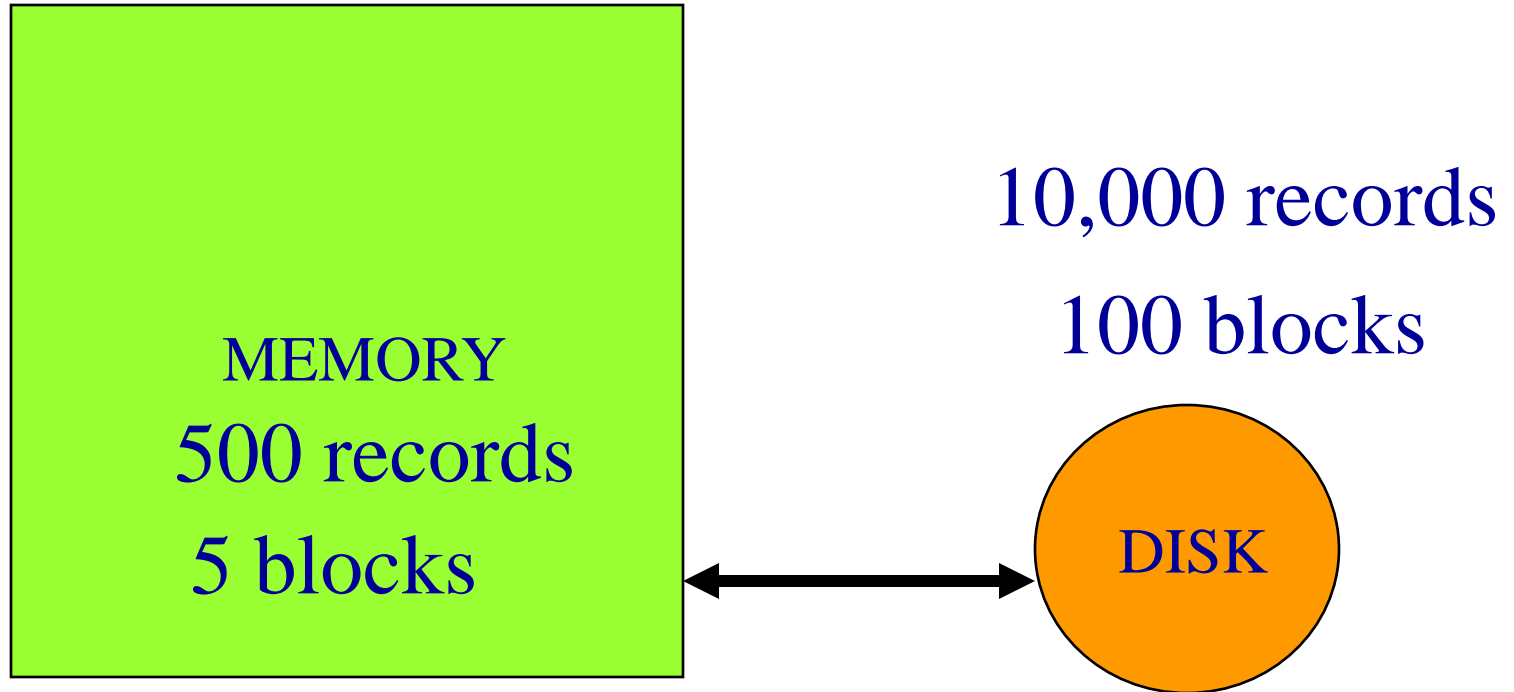
# External Merge Sort

- Sort 10,000 records.
- Enough memory for 500 records.
- Block size is 100 records.
- $t_{IO}$  = time to input/output 1 block  
(includes seek, latency, and transmission times)
- $t_{IS}$  = time to internally sort 1 memory load
- $t_{IM}$  = time to internally merge 1 block load

# External Merge Sort

- Two phases.
  - Run generation.
    - A run is a sorted sequence of records.
  - Run merging.

# Run Generation




- Input 5 blocks.
- Sort.
- Output as a run.
- Do 20 times.

- $5t_{IO}$
- $t_{IS}$
- $5t_{IO}$
- $200t_{IO} + 20t_{IS}$


# Run Merging

- Merge Pass.
  - Pairwise merge the 20 runs into 10.
  - In a merge pass all runs (except possibly one) are pairwise merged.
- Perform 4 more merge passes, reducing the number of runs to 1.

# Merge 20 Runs



# Merge R1 and R2



- Fill **I0** (Input 0) from **R1** and **I1** from **R2**.
- Merge from **I0** and **I1** to output buffer.
- Write whenever output buffer full.
- Read whenever input buffer empty.



# Time To Merge R1 and R2

- Each is 5 blocks long.
- Input time =  $10t_{IO}$ .
- Write/output time =  $10t_{IO}$ .
- Merge time =  $10t_{IM}$ .
- Total time =  $20t_{IO} + 10t_{IM}$ .

# Time For Pass 1 (R→S)

- Time to merge one pair of runs  
 $= 20t_{IO} + 10t_{IM}.$
- Time to merge all 10 pairs of runs  
 $= 200t_{IO} + 100t_{IM}.$

# Time To Merge S1 and S2

- Each is 10 blocks long.
- Input time =  $20t_{IO}$ .
- Write/output time =  $20t_{IO}$ .
- Merge time =  $20t_{IM}$ .
- Total time =  $40t_{IO} + 20t_{IM}$ .

## Time For Pass 2 (S $\rightarrow$ T)

- Time to merge one pair of runs  
 $= 40t_{IO} + 20t_{IM}.$
- Time to merge all 5 pairs of runs  
 $= 200t_{IO} + 100t_{IM}.$

# Time For One Merge Pass

- Time to input all blocks =  $100t_{IO}$ .
- Time to output all blocks =  $100t_{IO}$ .
- Time to merge all blocks =  $100t_{IM}$ .
- Total time for a merge pass =  $200t_{IO} + 100t_{IM}$ .

# Total Run-Merging Time


- (time for one merge pass) \* (number of passes)  
= (time for one merge pass)  
\*  $\text{ceil}(\log_2(\text{number of initial runs}))$   
=  $(200t_{\text{IO}} + 100t_{\text{IM}}) * \text{ceil}(\log_2(20))$   
=  $(200t_{\text{IO}} + 100t_{\text{IM}}) * 5$

# Factors In Overall Run Time

- Run generation.  $200t_{IO} + 20t_{IS}$ 
  - Internal sort time.
  - Input and output time.
- Run merging.  $(200t_{IO} + 100t_{IM}) * \text{ceil}(\log_2(20))$ 
  - Internal merge time.
  - Input and output time.
  - Number of initial runs.
  - Merge order (number of merge passes is determined by number of runs and merge order)

# Improve Run Generation

- Overlap input, output, and internal sorting.






# Improve Run Generation

- Generate runs whose length (on average) exceeds memory size.
- Equivalent to reducing number of runs generated.

# Improve Run Merging


- Overlap input, output, and internal merging.



# Improve Run Merging

- Reduce number of merge passes.
  - Use higher-order merge.
  - Number of passes  
=  $\text{ceil}(\log_k(\text{number of initial runs}))$   
where  $k$  is the merge order.

# Merge 20 Runs Using 5-Way Merging




Number of passes = 2

# I/O Time Per Merge Pass

- Number of input buffers needed is linear in merge order  $k$ .
- Since memory size is fixed, block size decreases as  $k$  increases (after a certain  $k$ ).
- So, number of blocks increases.
- So, number of seek and latency delays per pass increases.


# I/O Time Per Merge Pass



# Total I/O Time To Merge Runs



- (I/O time for one merge pass)  
\*  $\text{ceil}(\log_k(\text{number of initial runs}))$

↑  
Total  
I/O  
time to  
merge  
runs



merge order  $k$  →


# Internal Merge Time



- Naïve way  $\Rightarrow k - 1$  compares to determine next record to move to the output buffer.
- Time to merge  $n$  records is  $c(k - 1)n$ , where  $c$  is a constant.
- Merge time per pass is  $c(k - 1)n$ .
- Total merge time is  $c(k - 1)n \log_k r \sim cn(k/\log_2 k) \log_2 r$ .



# Merge Time Using A Selection Tree



- Time to merge  $n$  records is  $dn\log_2 k$ , where  $d$  is a constant.
- Merge time per pass is  $dn\log_2 k$ .
- Total merge time is  $(dn\log_2 k) \log_k r = dn\log_2 r$ .